

Machine Learning with the Experts: School Budgets

4). Learning from the experts:

a). Deciding what's a word

Import the CountVectorizer

```
from sklearn.feature_extraction.text import CountVectorizer
```

Create the text vector

```
text_vector = combine_text_columns(X_train)
```

Create the token pattern: TOKENS_ALPHANUMERIC

```
TOKENS_ALPHANUMERIC = '[A-Za-z0-9]+(?=\s+)'
```

Instantiate the CountVectorizer: text_features

```
text_features = CountVectorizer(token_pattern=TOKENS_ALPHANUMERIC)
```

Fit text_features to the text vector

```
text_features.fit(text_vector)
```

Print the first 10 tokens

```
print(text_features.get_feature_names()[:10])
```

<script.py> output:

```
['00a', '12', '1st', '2nd', '3rd', '5th', '70', '70h', '8', 'a']
```

b). N-gram range in sci-kit learn

```
# Import pipeline
```

```
from sklearn.pipeline import Pipeline
```

```
# Import classifiers
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.multiclass import OneVsRestClassifier
```

```
# Import CountVectorizer
```

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
# Import other preprocessing modules
```

```
from sklearn.preprocessing import Imputer
```

```
from sklearn.feature_selection import chi2, SelectKBest
```

```
# Select 300 best features
```

```
chi_k = 300
```

```
# Import functional utilities
```

```
from sklearn.preprocessing import FunctionTransformer, MaxAbsScaler
```

```
from sklearn.pipeline import FeatureUnion
```

```
# Perform preprocessing
```

```
get_text_data = FunctionTransformer(combine_text_columns, validate=False)
```

```
get_numeric_data = FunctionTransformer(lambda x: x[NUMERIC_COLUMNS], validate=False)
```

```
# Create the token pattern: TOKENS_ALPHANUMERIC
```

```
TOKENS_ALPHANUMERIC = '[A-Za-z0-9]+(?=\s+)'
```

```
# Instantiate pipeline: pl
```

```
pl = Pipeline([
```

```
('union', FeatureUnion(  
    transformer_list = [  
        ('numeric_features', Pipeline(  
            ('selector', get_numeric_data),  
            ('imputer', Imputer())  
        )),  
        ('text_features', Pipeline(  
            ('selector', get_text_data),  
            ('vectorizer', CountVectorizer(token_pattern=TOKENS_ALPHANUMERIC,  
                                           ngram_range=(1,2))),  
            ('dim_red', SelectKBest(chi2, chi_k)  
        ))  
    ]  
)),  
('scale', MaxAbsScaler()),  
('clf', OneVsRestClassifier(LogisticRegression()))  
])
```

c). Implement Interaction modeling in scikit learn**# Instantiate pipeline: pl**

```
pl = Pipeline([
    ('union', FeatureUnion(
        transformer_list = [
            ('numeric_features', Pipeline([
                ('selector', get_numeric_data),
                ('imputer', Imputer())
            ])),
            ('text_features', Pipeline([
                ('selector', get_text_data),
                ('vectorizer', CountVectorizer(token_pattern=TOKENS_ALPHANUMERIC,
                                                ngram_range=(1, 2))),
                ('dim_red', SelectKBest(chi2, chi_k))
            ]))
        ]
    )),
    ('int', SparseInteractions(degree=2)),
    ('scale', MaxAbsScaler()),
    ('clf', OneVsRestClassifier(LogisticRegression()))
])
```

d). Implementing the hashing trick in scikit-learn**# Import HashingVectorizer****from sklearn.feature_extraction.text import HashingVectorizer****# Get text data: text_data****text_data = combine_text_columns(X_train)****# Create the token pattern: TOKENS_ALPHANUMERIC****TOKENS_ALPHANUMERIC = '[A-Za-z0-9]+(?=\s+)'****# Instantiate the HashingVectorizer: hashing_vec****hashing_vec = HashingVectorizer(token_pattern=TOKENS_ALPHANUMERIC)****# Fit and transform the Hashing Vectorizer****hashed_text = hashing_vec.fit_transform(text_data)****# Create DataFrame and print the head****hashed_df = pd.DataFrame(hashed_text.data)****print(hashed_df.head())**

<script.py> output:

```
0
0 -0.160128
1 0.160128
2 -0.480384
3 -0.320256
4 0.160128
```

e). Build the winning model**# Import the hashing vectorizer****from sklearn.feature_extraction.text import HashingVectorizer****# Instantiate the winning model pipeline: pl**

```
pl = Pipeline([  
    ('union', FeatureUnion(  
        transformer_list = [  
            ('numeric_features', Pipeline([  
                ('selector', get_numeric_data),  
                ('imputer', Imputer()  
            )]),  
            ('text_features', Pipeline([  
                ('selector', get_text_data),  
                ('vectorizer', HashingVectorizer(token_pattern=TOKENS_ALPHANUMERIC,  
                    non_negative=True, norm=None, binary=False,  
                    ngram_range=(1,2))),  
                ('dim_red', SelectKBest(chi2, chi_k))  
            ]))  
        ]  
    )),  
    ('int', SparseInteractions(degree=2)),  
    ('scale', MaxAbsScaler()),  
    ('clf', OneVsRestClassifier(LogisticRegression()))  
])
```